

1 Maximum Likelihood Estimation

Maximum Likelihood Estimation (MLE) is a method of estimating the parameters of a statistical model given observations, by finding the parameters that maximize the likelihood of the observations. Concretely, given observations y_1, y_2, \dots, y_n distributed according to $p_\theta(y_1, y_2, \dots, y_n)$ (here p_θ can be a probability mass function for discrete observations or a density for continuous observations), the likelihood function is defined as $L(\theta) = p_\theta(y_1, y_2, \dots, y_n)$ and the MLE is

$$\hat{\theta}_{\text{MLE}} = \arg \max_{\theta} L(\theta).$$

We often make the assumption that the observations are *independent and identically distributed* or iid, in which case $p_\theta(y_1, y_2, \dots, y_n) = p_\theta(y_1) \cdot p_\theta(y_2) \cdot \dots \cdot p_\theta(y_n)$.

- (a) Your friendly TA recommends maximizing the log-likelihood $\ell(\theta) = \log L(\theta)$ instead of $L(\theta)$. Why does this yield the same solution $\hat{\theta}_{\text{MLE}}$? Why is it easier to solve the optimization problem for $\ell(\theta)$ in the iid case? Write down both $L(\theta)$ and $\ell(\theta)$ for the Gaussian $f_\theta(y) = \frac{1}{\sqrt{2\pi\sigma^2}} e^{\frac{-(y-\mu)^2}{2\sigma^2}}$ with $\theta = (\mu, \sigma)$.

(b) What is $\int p_\theta(y_1, y_2, \dots, y_n) dy_1 \cdots dy_n$? Can we say anything about $\int p_\theta(y_1, y_2, \dots, y_n) d\theta$?

(c) The Poisson distribution is $f_\lambda(y) = \frac{\lambda^y e^{-\lambda}}{y!}$. Let Y_1, Y_2, \dots, Y_n be a set of independent and identically distributed random variables with Poisson distribution with parameter λ . Find the joint distribution of Y_1, Y_2, \dots, Y_n . Find the maximum likelihood estimator of λ as a function of observations y_1, y_2, \dots, y_n .

2 ℓ_1 - and ℓ_2 -Regularization

Consider sample points $X_1, X_2, \dots, X_n \in \mathbb{R}^d$ and associated values $y_1, y_2, \dots, y_n \in \mathbb{R}$, an $n \times d$ design matrix $X = [X_1 \ \dots \ X_n]^\top$ and an n -vector $y = [y_1 \ \dots \ y_n]^\top$.

For the sake of simplicity, assume (1) that the sample data have been centered (i.e each feature has mean 0) and (2) that the sample data have been whitened, meaning a linear transformation of is applied to the original data matrix so that the resulting features have variance 1 and the features are uncorrelated; i.e., $X^\top X = nI$.

For this question, we will not use a fictitious dimension nor a bias term; our linear regression function will output zero for $x = 0$.

Consider linear least-squares regression with regularization in the ℓ_1 -norm, also known as Lasso. The Lasso cost function is

$$J(w) = \|Xw - y\|^2 + \lambda \|w\|_1$$

where $w \in \mathbb{R}^d$ and $\lambda > 0$ is the regularization parameter. Let $w^* = \arg \min_{w \in \mathbb{R}^d} J(w)$ denote the weights that minimize the cost function.

In the following steps, we will explore the sparsity-promoting property of the ℓ_1 -norm and compare this with the ℓ_2 -norm.

1. We use the notation X_{*i} to denote column i of the design matrix X , which represents the i^{th} feature. Write $J(w)$ in the following form for appropriate functions g and f .

$$J(w) = g(y) + \sum_{i=1}^d f(X_{*i}, w_i, y, \lambda)$$

2. If $w_i^* > 0$, what is the value of w_i^* ?

3. If $w_i^* < 0$, what is the value of w_i^* ?

4. Considering parts 2 and 3, what is the condition for w_i^* to be zero?

5. Now consider ridge regression, which uses the ℓ_2 regularization term $\lambda |w|^2$. How does this change the function $f(\cdot)$ from part 1? What is the new condition in which $w_i^* = 0$? How does

it differ from the condition you obtained in part 4?

3 Probabilistic Interpretation of Lasso

Let's start with the probabilistic interpretation of least squares. Start with labels $y \in \mathbb{R}$, data $\mathbf{x} \in \mathbb{R}^d$, and noise $z \sim \mathcal{N}(0, \sigma^2)$, where $y = \mathbf{w}^T \mathbf{x} + z$. Recall from a previous homework that we then have

$$P(y|\mathbf{x}, \sigma^2) \sim \mathcal{N}(\mathbf{w}^T \mathbf{x}, \sigma^2)$$

However, maximum likelihood estimates (MLE) can overfit by picking parameters that mirror the training data. To ameliorate this issue, we can assume a Laplace prior on $w_j \sim \text{Laplace}(0, t)$, i.e.

$$P(w_j) = \frac{1}{2t} e^{-|w_j|/t}$$

$$P(\mathbf{w}) = \prod_{j=1}^D P(w_j) = \left(\frac{1}{2t}\right)^D \cdot e^{-\frac{\sum |w_j|}{t}}$$

Here, we will see that this modification results in a new objective function, called Lasso.

Recall that the MLE objective finds the parameters that maximize the likelihood of the data,

$$\begin{aligned} \mathbf{w}^* &= \arg \max_{\mathbf{w}} L(\mathbf{w}) \\ &= \arg \max_{\mathbf{w}} P(Y_1, \dots, Y_n, |\mathbf{w}, \mathbf{X}_1, \dots, \mathbf{X}_n, \sigma^2) \\ &= \arg \max_{\mathbf{w}} \prod_{i=1}^n P(Y_i | \mathbf{X}_i, \mathbf{w}, \sigma^2). \end{aligned}$$

When working in a Bayesian framework, we instead focus on the posterior distribution of the parameters conditioned on the data, $P(\mathbf{w} | Y_1, \dots, Y_n, \mathbf{X}_1, \dots, \mathbf{X}_n, \sigma^2)$. To pick a single model, we can choose the \mathbf{w} that is most likely according to the posterior,

$$\begin{aligned} \mathbf{w}^* &= \arg \max_{\mathbf{w}} P(\mathbf{w} | Y_1, \dots, Y_n, \mathbf{X}_1, \dots, \mathbf{X}_n, \sigma^2) \\ &= \arg \max_{\mathbf{w}} \frac{P(\mathbf{w}, Y_1, \dots, Y_n | \mathbf{X}_1, \dots, \mathbf{X}_n, \sigma^2)}{P(Y_1, \dots, Y_n | \mathbf{X}_1, \dots, \mathbf{X}_n, \sigma^2)} \\ &= \arg \max_{\mathbf{w}} \frac{P(Y_1, \dots, Y_n, |\mathbf{w}, \mathbf{X}_1, \dots, \mathbf{X}_n, \sigma^2) P(\mathbf{w})}{P(Y_1, \dots, Y_n | \mathbf{X}_1, \dots, \mathbf{X}_n, \sigma^2)} \\ &= \arg \max_{\mathbf{w}} \frac{L(\mathbf{w}) P(\mathbf{w})}{P(Y_1, \dots, Y_n)} \\ &= \arg \max_{\mathbf{w}} L(\mathbf{w}) P(\mathbf{w}) \quad \text{since } P(Y_1, \dots, Y_n | \mathbf{X}_1, \dots, \mathbf{X}_n, \sigma^2) \text{ does not depend on } \mathbf{w}. \end{aligned}$$

We call \mathbf{w}^* the Maximum a posteriori (MAP) estimate.

(a) Write the log-likelihood for this MAP estimate.

(b) We already have the log-likelihood for MAP. Show that MAP—in this case, Gaussian noise with a Laplace prior—is equivalent to minimizing the following. Additionally, identify the constant λ . Note that $\|\mathbf{w}\|_1 = \sum_{j=1}^D |w_j|$.

$$J(\mathbf{w}) = \sum_{i=1}^n (Y_i - \mathbf{w}^T \mathbf{X}_i)^2 + \lambda \|\mathbf{w}\|_1$$