

1 Logistic Regression

Assume that we have n i.i.d. data points $(\mathbf{x}_1, y_1), \dots, (\mathbf{x}_n, y_n)$, where each y_i is a binary label in $\{0, 1\}$. We model the posterior probability as a Bernoulli distribution and the probability for each class is the sigmoid function, i.e., $p(y|\mathbf{x}; \mathbf{w}) = q^y(1 - q)^{1-y}$, where $q = s(\mathbf{w}^\top \mathbf{x})$ and $s(\zeta) = \frac{1}{1+e^{-\zeta}}$ is the sigmoid function. Write out the likelihood and log likelihood functions. Comment on whether it is possible to find a closed form maximum likelihood estimate of \mathbf{w} , and describe an alternate approach.

Solution:

The likelihood is:

$$L(\mathbf{w}) = \prod_{i=1}^n p(y = y_i | \mathbf{x}_i) = \prod_{i=1}^n q_i^{y_i} (1 - q_i)^{1-y_i}.$$

Now maximizing the likelihood of the training data as a function of the parameters \mathbf{w} , we get:

$$\begin{aligned} \hat{\mathbf{w}} &= \arg \max_{\mathbf{w}} L(\mathbf{w}) = \arg \max_{\mathbf{w}} \prod_{i=1}^n q_i^{y_i} (1 - q_i)^{1-y_i} \\ &= \arg \max_{\mathbf{w}} \sum_{i=1}^n y_i \log(q_i) + (1 - y_i) \log(1 - q_i) \\ &= \arg \max_{\mathbf{w}} \sum_{i=1}^n y_i \log\left(\frac{q_i}{1 - q_i}\right) + \log(1 - q_i) \end{aligned}$$

Since q_i is the sigmoid function, we get that:

$$\begin{aligned} \hat{\mathbf{w}} &= \arg \max_{\mathbf{w}} \sum_{i=1}^n y_i \log\left(\frac{q_i}{1 - q_i}\right) + \log(1 - q_i) \\ &= \arg \max_{\mathbf{w}} \sum_{i=1}^n y_i \log\left(\frac{\frac{1}{1+\exp\{-\mathbf{w}^\top \mathbf{x}_i\}}}{1 - \frac{1}{1+\exp\{-\mathbf{w}^\top \mathbf{x}_i\}}}\right) + \log\left(1 - \frac{1}{1 + \exp\{-\mathbf{w}^\top \mathbf{x}_i\}}\right) \\ &= \arg \max_{\mathbf{w}} \sum_{i=1}^n y_i \log\left(\frac{\frac{1}{1+\exp\{\mathbf{w}^\top \mathbf{x}_i\}}}{\frac{1+\exp\{-\mathbf{w}^\top \mathbf{x}_i\}-1}{1+\exp\{-\mathbf{w}^\top \mathbf{x}_i\}}}\right) + \log\left(\frac{1 + \exp\{-\mathbf{w}^\top \mathbf{x}_i\} - 1}{1 + \exp\{-\mathbf{w}^\top \mathbf{x}_i\}}\right) \\ &= \arg \max_{\mathbf{w}} \sum_{i=1}^n y_i \log\left(\frac{1}{\exp\{-\mathbf{w}^\top \mathbf{x}_i\}}\right) + \log\left(\frac{\exp\{-\mathbf{w}^\top \mathbf{x}_i\}}{1 + \exp\{-\mathbf{w}^\top \mathbf{x}_i\}}\right) \end{aligned}$$

$$= \arg \max_{\mathbf{w}} \sum_{i=1}^n y_i \mathbf{w}^\top \mathbf{x}_i - \mathbf{w}^\top \mathbf{x}_i - \log(1 + \exp\{-\mathbf{w}^\top \mathbf{x}_i\})$$

Multiplying everything by -1 and switching to minimization we have:

$$\hat{\mathbf{w}} = - \arg \min_{\mathbf{w}} \sum_{i=1}^n (1 - y_i) \mathbf{w}^\top \mathbf{x}_i + \log(1 + \exp\{-\mathbf{w}^\top \mathbf{x}_i\})$$

Let us denote $J(\mathbf{w}) = - \sum_{i=1}^n (1 - y_i) \mathbf{w}^\top \mathbf{x}_i - \log(1 + \exp\{-\mathbf{w}^\top \mathbf{x}_i\})$. Notice that $J(\mathbf{w})$ is convex in \mathbf{w} , so global minima can be found. Recall that $s'(\zeta) = s(\zeta)(1 - s(\zeta))$. Now let us take the derivative of $J(\mathbf{w})$ w.r.t \mathbf{w} :

$$\frac{\partial J}{\partial \mathbf{w}} = - \sum_{i=1}^n (1 - y_i) \mathbf{x}_i - \frac{\exp\{-\mathbf{w}^\top \mathbf{x}_i\}}{1 + \exp\{-\mathbf{w}^\top \mathbf{x}_i\}} \mathbf{x}_i = \sum_{i=1}^n (1 - s(\mathbf{w}^\top \mathbf{x}_i) - y_i + 1) \mathbf{x}_i = \sum_{i=1}^n (s_i - y_i) \mathbf{x}_i = \mathbf{X}^\top (\mathbf{s} - \mathbf{y})$$

where, $s_i = s(-\mathbf{w}^\top \mathbf{x}_i)$, $\mathbf{s} = (s_1, \dots, s_n)^\top$, $\mathbf{y} = (y_1, \dots, y_n)^\top$ and $\mathbf{X} = \begin{bmatrix} \mathbf{x}_1^\top \\ \vdots \\ \mathbf{x}_n^\top \end{bmatrix}$.

Unfortunately, we can't get a closed form estimate for \mathbf{w} by setting the derivative to zero.

2 Initialization of Weights for Backpropagation

Assume a fully-connected 1-hidden-layer network. Denote the dimensionalities of the input, hidden, and output layers as $d^{(0)}$, $d^{(1)}$, and $d^{(2)}$. That is, the input (which we will denote with a superscript (0)) has dimensions $x_1^{(0)}, \dots, x_{d^{(0)}}^{(0)}$. Let g denote the activation function applied at each layer. As defined in lecture, let $S_j^{(l)} = \sum_{i=1}^{d^{(l-1)}} w_{ij}^{(l)} x_i^{(l-1)}$ be the weighted input to node j in layer l , and let $\delta_j^{(l)} = \frac{\partial \ell}{\partial S_j^{(l)}}$ be the partial derivative of the final loss ℓ with respect to $S_j^{(l)}$.

Recall that backpropagation is simply an efficient method to compute the gradient of the loss function so we can use it with gradient descent. These methods require the parameters to be initialized to some value. In logistic regression we were able to initialize all weights as 0.

- (a) Imagine that we initialize the values of our weights to be some constant w . After performing the forward pass, what is the value of $x_j^{(1)}$ in terms of the elements of $\{x_i^{(0)} : i = 1, \dots, d^{(0)}\}$? What is the relationship between each $x_j^{(1)}$?

Solution: Since all of the weights are equal, we have:

$$x_j^{(1)} = g\left(\sum_{i=1}^{d^{(0)}} w_{ij}^{(1)} x_i^{(0)}\right) = g\left(w \sum_{i=1}^{d^{(0)}} x_i^{(0)}\right)$$

Note that this equation does not depend on j so all $x_j^{(1)}$ are equal.

- (b) After the backward pass of backpropagation, what is the relation between the members of the set $\{\delta_i^{(1)} : i = 1, \dots, d^{(1)}\}$, assuming we have calculated $\{\delta_j^{(2)} : j = 1, \dots, d^{(2)}\}$?

Solution: Since all of the weights are equal:

$$\begin{aligned} \delta_i^{(1)} &= \sum_{j=1}^{d^{(2)}} \delta_j^{(2)} w_{ij}^{(1)} g'(S_i^{(1)}) \\ &= g'(S_i^{(1)}) w \sum_{j=1}^{d^{(2)}} \delta_j^{(2)} \\ &= g'\left(w \sum_{k=1}^{d^{(0)}} x_k^{(0)}\right) w \sum_{j=1}^{d^{(2)}} \delta_j^{(2)} \end{aligned}$$

The sum term is the same for all values of i in layer 1. The members of the set are equal.

- (c) For a reasonable loss function, can we say the same about each $\delta_i^{(2)}$?

Solution: No. $\delta_i^{(2)}$ depends on y_i , which is different for each i . Note that *any* reasonable output loss/activation should result in $\delta_i^{(2)}$ depending on a target y_i (otherwise the output is not attempting to approach a particular target value).

- (d) After the weights are updated and one iteration of gradient descent has been completed, what can we say about the weights?

Solution: Our gradient descent update looks like this:

$$\begin{aligned} w_{ij}^{(l)} &= w_{ij}^{(l)} - \eta \delta_j^{(l)} x_i^{(l-1)} = w - \eta \delta_j^{(l)} x_i^{(l-1)} \\ \implies w_{ij}^{(1)} &= w - \eta \delta_j^{(1)} x_i^{(0)} \\ \implies w_{ij}^{(2)} &= w - \eta \delta_j^{(2)} x_i^{(1)} \end{aligned}$$

For a given i , $w_{ij}^{(1)}$ will be the same for all j because $\delta_j^{(1)}$ is equal for all j .

For a given j , $w_{ij}^{(2)}$ will be the same for all i because $x_i^{(1)}$ is equal for all i .

- (e) Even though $w_{ij}^{(2)}$ is different for each j (but for a fixed j , it is the same for each i), this pattern continues. Why?

Solution: Let $w_{ij}^{(1)} = w_i^{(1)}$ for all i . Let $w_{ij}^{(2)} = w_j^{(2)}$ for all j . Then:

$$x_j^{(1)} = g \left(\sum_{i=1}^{d^{(0)}} w_{ij}^{(1)} x_i^{(0)} \right) = g \left(\sum_{i=1}^{d^{(0)}} w_i^{(1)} x_i^{(0)} \right)$$

Which does not depend on j . Also,

$$\delta_i^{(1)} = \sum_{j=1}^{d^{(2)}} \delta_j^{(2)} w_{ij}^{(2)} g' \left(\sum_{k=1}^{d^{(0)}} w_{ki}^{(1)} x_k^{(0)} \right) = \sum_{j=1}^{d^{(2)}} \delta_j^{(2)} w_j^{(2)} g' \left(\sum_{k=1}^{d^{(0)}} w_k^{(1)} x_k^{(0)} \right)$$

Which does not depend on j .

All $x_j^{(1)}$ being the same and all $\delta_i^{(1)}$ being the same will continue no matter how many iterations you perform. Intuitively, $x_j^{(1)}$ will always be the same for all j , due to the “outgoing” symmetry in the weights in layer 1, and the $\delta_i^{(1)}$ will always be the same due to the “incoming” symmetry in the weights in layer 2.

- (f) To solve this problem, we randomly initialize our weights. This is called symmetry breaking. Why are we able to set our weights to 0 for logistic regression?

Solution: Logistic regression is a fully-connected neural network with one output node and zero hidden layers (remember that the $\delta_j^{(2)}$'s are different).

Another reason: in logistic regression, the loss function is convex. Any starting point should lead us to the global optimum.