

1 Maximum Likelihood Review

Suppose you are collecting data on the relative rates of different types of twins, and you obtain the following observations:

- there are m_i pairs of identical male twins and f_i pairs of identical female twins
- there are m_f pairs of fraternal (not identical) male twins and f_f pairs of fraternal female twins
- there are b pairs of fraternal opposite gender twins

To model this data, we choose these distributions and parameters:

- Given that a pair of siblings are twins, they are identical with probability θ , and fraternal (non-identical) with probability $1 - \theta$
- Given they are identical twins, the twins are both male with probability p and both female with probability $1 - p$.
- Given they are twins and not identical (and thus are fraternal twins), the probability of both male twins is q^2 , probability of both female twins is $(1 - q)^2$ and probability of opposite gender twins is $2q(1 - q)$.

- (a) Write expressions for the likelihood and the log-likelihood of the data as functions of the parameters θ , p , and q for the observations m_i , f_i , m_f , f_f , b .

Likelihood $L(\theta, p, q) =$

Log likelihood $l(\theta, p, q) =$

- (b) What are the maximum likelihood estimates for θ , p and q ? Scratch space is provided to you here, which you may find useful.

2 MAP Estimation Review

Suppose we have a data set of n data points $D = \{x_1, \dots, x_n\}$, with each point drawn independently from a Gaussian with mean μ and variance σ^2 .

$$p(x) = \frac{1}{\sigma \sqrt{2\pi}} \exp\left(-\frac{(x - \mu)^2}{2\sigma^2}\right)$$

We will place the following prior on μ :

$$\mu \sim \mathcal{N}(\mu_0, \sigma_0^2)$$

The MAP estimate of μ is defined as

$$\mu_{\text{MAP}} = \arg \max_{\mu} p(\mu|D)$$

(a) Write an expression for the MAP estimate of μ .

(b) What happens to the MAP estimate as $\sigma_0^2 \rightarrow \infty$, and how is this estimate related to the ML estimate? Interpret this result.

(c) What happens to the MAP estimate as $\sigma^2 \rightarrow \infty$?

3 Prediction Error of Ridge Regression

- (a) Let A be a $d \times n$ matrix and B be a $n \times d$ matrix. For any $\mu > 0$, show that $(AB + \mu I)^{-1}A = A(BA + \mu I)^{-1}$, if $AB + \mu I$ and $BA + \mu I$ are invertible.

- (b) Let $X \in \mathbb{R}^{n \times d}$ be n samples of d features, and $y \in \mathbb{R}^n$ be the corresponding n samples of the quantity that you would like to predict with regression. Let

$$\widehat{\theta}_\lambda = \arg \min_{\theta} \|X\theta - y\|_2^2 + \lambda \|\theta\|_2^2,$$

for $\lambda > 0$, be the solution to the ridge regression problem.

Using part (a), show that $\widehat{\theta}_\lambda = X^\top (XX^\top + \lambda I)^{-1}y$.

- (c) Suppose X has the singular value decomposition $U\Sigma V^\top$, where $\Sigma = \text{diag}(s_1, \dots, s_d)$, $s_i \geq 0$. Show that $\widehat{\theta}_\lambda = VDU^\top y$, where D is a diagonal matrix to be determined.

- (d) Let $\widehat{y}_\lambda = X\widehat{\theta}_\lambda$ be the predictions made by the ridge regressor $\widehat{\theta}_\lambda$. Suppose we have $y = X\theta_* + z$, where $\theta_* \in \mathbb{R}^d$ and $z = \mathcal{N}(0, \sigma^2 I) \in \mathbb{R}^n$ ($\sigma > 0$). Further suppose that X is an orthogonal matrix, that is, $X^\top X = I$.

$\mathbb{E}\|X(\widehat{\theta}_\lambda - \theta_*)\|^2$ is the expected squared difference between the predictions made by the ridge regressor \widehat{y}_λ and $X\theta_*$, where the expectation is taken with respect to z ($\|\cdot\|$ denotes ℓ_2 norm).

Show that $\mathbb{E}\|X(\widehat{\theta}_\lambda - \theta_*)\|^2 = \frac{1}{(1+\lambda)^2} (\lambda^2 \|\theta_*\|^2 + d\sigma^2)$.

- (e) What is the λ^* that you should pick to minimize the prediction error you computed in part (d)? Comment on how d , σ^2 , and θ_* affect the optimal choice of the regularization parameter λ .