

1 Learning about PyTorch

PyTorch is a Python library that is widely used for deep learning. PyTorch greatly simplifies the process of building deep learning models, training them via backpropagation and stochastic gradients, loading and processing data, and more. You will need to gain a working understanding of PyTorch in order to successfully complete the homework, and this understanding will quite possibly benefit you later on in your future machine learning endeavors.

For this discussion, you will use https://pytorch.org/tutorials/beginner/deep_learning_60min_blitz.html as a reference to answer the questions below, rather than lecture materials. If you are not yet familiar with PyTorch, it is strongly encouraged that you review this excellent tutorial in detail before diving into the homework assignment. For the purposes of today's discussion, do your best to skim through the tutorial and search for the answers to the questions. Some of these answers will not be directly provided in the text and will instead require you to read the code and infer the behavior of certain functions.

Solution: Check out <https://tinyurl.com/PyTorch189> for TA Sean Lin's interactive Colab tutorial.

- (a) Compare and contrast NumPy's `ndarray` with PyTorch's `tensor`. In what ways are these two types similar? In what ways are they different?

Solution: Both `ndarray` and `tensor` represent multidimensional arrays. The difference between a NumPy `ndarray` and a PyTorch `tensor` is that `tensors` can be backed by the accelerator memory, such as a GPU, and they store numerical gradient information in a `grad` attribute when instantiated with `requires_grad=True`.

On setting `requires_grad=True`, `tensors` start forming a backward graph that tracks every operation applied on them to calculate the gradients using something called a dynamic computation graph (DCG). By PyTorch's design, gradients can only be calculated for floating point tensors.

- (b) What is the purpose of the `torch.nn.Module` class? What functionality does it provide?

Solution: `torch.nn.Module` is the base class for all neural network modules. Your custom models and layers should also subclass this class.

`Modules` can contain other `Modules`, allowing to nest them in a tree structure. Inheriting from `torch.nn.Module` provides useful functionality to your class. For example, it makes your class keep track of its trainable parameters via the `parameters` bound method. You can also

easily swap a `Module` between CPU and GPU with the `to(device)` bound method, where `device` can be a CPU (`torch.device("cpu")`) or GPU (`torch.device("cuda:0")`).

- (c) What function is required to make a `Module` callable? What does calling the `Module` do for neural networks?

Solution: The `forward` function is required. We just have to define the `forward` function, and the backward function (where gradients are computed) is automatically defined for you by PyTorch. We can use any of the Tensor operations in the `forward` function.

Note that, just like in Homework 2, certain intermediate values are stored during the forward pass of a neural network which will be used in backpropagation, and this is automatically handled by the built-in PyTorch layers. Using the `torch.no_grad` context manager disables this functionality and can make evaluation more efficient.

When calling the `Module`, the `__call__` function is already defined in `torch.nn.Module`. You should call the `Module` directly (`output = model(data)`) instead of something like `output = model.forward(data)`, because calling the `Module` performs additional functionality such as registering hooks.

- (d) What PyTorch function is used when moving between convolution layers and fully connected layers in a neural network forward pass?

Solution: The `flatten` function is used to transform the activations of the convolution layers, which have channel, height, and width dimensions, into flat vectors that are input into the fully connected layers.

- (e) How is backpropagation performed in PyTorch?

Solution: At an interface level, it is relatively simple. A `loss`, which is a `tensor` is computed by comparing the outputs of the model with the desired outputs (labels), using a `criterion` such as `torch.nn.CrossEntropyLoss`. This `loss` has a `backward` bound method which can be invoked that automatically populates the `grad` variables of all `tensors` with `requires_grad=True` that were involved in the computation of `loss`.

More details:

PyTorch's automatic differentiation engine (`torch.autograd`) calculates derivatives. It records a graph of all the operations performed on a gradient enabled tensor and creates an acyclic graph called the dynamic computational graph. The leaves of this graph are input tensors and the roots are output tensors. Gradients are calculated by tracing the graph from the root to the leaf and multiplying every gradient in the way using the chain rule. This is effectively an automatic and efficient implementation of the forward and backward pass you implemented in Homework 2.

PyTorch builds a Dynamic Computational Graph (DCG) from scratch in every iteration providing maximum flexibility to gradient calculation. For example, for a forward operation (function) `Mul` a backward operation (function) called `MulBackward` is dynamically integrated in the backward graph for computing the gradient. Gradient enabled tensors (variables) along with functions (operations) combine to create the dynamic computational graph. The flow of

data and the operations applied to the data are defined at runtime hence constructing the computational graph dynamically. This graph is made dynamically by the `autograd` class under the hood. You don't have to encode all possible paths before you launch the training — what you run is what you differentiate. This property allows you to add or remove layers, or shuffle parameters between iterations.

`backward` is the method which actually calculates the gradient by passing its argument through the backward graph all the way up to every leaf node traceable from the calling root tensor. The calculated gradients are then stored in the `grad` attribute of every leaf node. Remember, the backward graph is already made dynamically during the forward pass. The `backward` method only calculates the gradient using the already made graph and stores them in leaf nodes.

- (f) What is the purpose of the `zero_grad` method? What happens if this method is not called?

Solution: In PyTorch, for every mini-batch during the training phase, we typically want to explicitly set the gradients to zero before backpropagation, because PyTorch accumulates the gradients on subsequent backward passes. The default action in PyTorch is to accumulate (i.e. sum) the gradients on every `loss.backward()` call.

Because of this, when we start a training loop, we should zero out the gradients so that we do the parameter update correctly. Otherwise, the gradient would be a combination of the old gradient, which we have already used to update the model parameters, and the newly-computed gradient.

- (g) How are model updates (e.g., via stochastic gradients) performed in PyTorch?

Solution: Models are updated by passing the `model.parameters()` into an optimizer, which typically comes from the `torch.optim` library (e.g., `torch.optim.SGD` or `torch.optim.Adam`). First, calling `loss.backward()` calculates and populates gradients with respect to the model parameters. Then, calling the optimizer's `step` bound method updates the parameters. Note that the optimizer does not update all tensors with a `grad` attribute, only the ones passed to it during initialization. An examples is shown in: https://pytorch.org/tutorials/beginner/blitz/autograd_tutorial.html#usage-in-pytorch

For more information see: https://pytorch.org/tutorials/beginner/basics/optimization_tutorial.html

- (h) What PyTorch class is responsible for processing and providing batches of data for training and testing? Briefly describe how it is constructed and used.

Solution: The `torch.utils.data.DataLoader` class.

Code for processing data samples can get messy and hard to maintain; we ideally want our dataset code to be decoupled from our model training code for better readability and modularity. PyTorch provides two data primitives: `torch.utils.data.DataLoader` and `torch.utils.data.Dataset`. These classes allow you to use pre-loaded datasets as well as your own data. `Dataset` stores the samples and their corresponding labels, and `DataLoader` wraps an iterable around the `Dataset` to enable easy access to the samples.

For more information see: https://pytorch.org/tutorials/beginner/basics/data_tutorial.html

- (i) How does one take advantage of GPUs when training and testing PyTorch models?

Solution: To make use of a GPU, we need to transfer the relevant `tensors` and `Modules` to the GPU first by calling `.cuda()` or `.to(device)`, where `device` corresponds to a GPU. Once all of the relevant `tensors` and `Modules` are on the GPU, the operations will automatically be computed on the GPU. To see an example see:

https://pytorch.org/tutorials/beginner/blitz/tensor_tutorial.html#tensor-operati

To additionally make sure of more than one GPU at once, we can make use of data parallelism or model parallelism on GPUs. Data parallelism is when you use the same model for every thread, but feed it with different parts of the data; model parallelism is when you use the same data for every thread, but split the model among threads.