

Due 10/11 at 11:59pm

- We prefer that you typeset your answers using \LaTeX or other word processing software. If you haven't yet learned \LaTeX , one of the crown jewels of computer science, now is a good time! Neatly handwritten and scanned solutions will also be accepted for the written questions.
- In all of the questions, **show your work**, not just the final answer.

Deliverables:

1. Submit a PDF of your homework to the Gradescope assignment entitled "HW3 Write-Up". **Please start each question on a new page.** If there are graphs, include those graphs in the correct sections. **Do not** put them in an appendix. We need each solution to be self-contained on pages of its own.
 - In your write-up, please state with whom you worked on the homework. This should be on its own page and should be the first page that you submit.
 - In your write-up, please copy the following statement and sign your signature next to it. (Mac Preview and FoxIt PDF Reader, among others, have tools to let you sign a PDF file.) We want to make it *extra* clear so that no one inadvertently cheats. *"I certify that all solutions are entirely in my own words and that I have not looked at another student's solutions. I have given credit to all external sources I consulted."*

1 Poisson Classification

Recall that the PDF of a Poisson random variable is

$$P(X = x) = e^{-\lambda} \frac{\lambda^x}{x!} \quad x \in \{0, 1, \dots, \infty\}$$

The PDF is defined for non-negative integral values.

You are given two classes ω_1, ω_2 of Poisson data with parameters λ_1 and λ_2 . This means that $x|\omega_1 \sim \text{Poisson}(\lambda_1)$ and $x|\omega_2 \sim \text{Poisson}(\lambda_2)$. Assume that $P(\omega_1) = P(\omega_2) = \frac{1}{2}$.

- (a) Find $P(\omega_1|x)$ in terms of λ_1 and λ_2 . What type of function is the posterior?
- (b) Find the optimal rule (decision boundary) for allocating an observation x to a particular class. In the case where $P(\omega_1) = P(\omega_2) = \frac{1}{2}$ and $\lambda_1 = 10$ and $\lambda_2 = 15$, calculate the decision boundary, probability of correct classification for each class, and total error rate.
- (c) Suppose instead of one, we can obtain two independent measurements x_1 and x_2 for the object to be classified. How does the allocation rule change? In the case where $P(\omega_1) = P(\omega_2) = \frac{1}{2}$ and $\lambda_1 = 10$ and $\lambda_2 = 15$, calculate the new total error.

2 Logistic posterior with exponential class conditionals

We have seen in class that Gaussian class conditionals can lead to a logistic posterior that is linear in X . Now, suppose the class conditionals are exponentially distributed with parameters λ_i , i.e.

$$p(x|Y = i) = \lambda_i \exp(-\lambda_i x), \quad \text{where } i \in \{0, 1\}$$

$$Y \sim \text{Bernoulli}(\pi)$$

Show that the posterior distribution of the class label given X is also a logistic function, however with a linear argument in X . What is the decision boundary?

3 Gaussian Classification

Let $f(x | C_i) \sim \mathcal{N}(\mu_i, \sigma^2)$ for a two-class, one-dimensional classification problem with classes C_1 and C_2 , $P(C_1) = P(C_2) = 1/2$, and $\mu_2 > \mu_1$.

- (a) Find the Bayes optimal decision boundary and the corresponding Bayes decision rule.
- (b) The Bayes error is the probability of misclassification,

$$P_e = P(\text{misclassified as } C_1 | C_2) P(C_2) + P(\text{misclassified as } C_2 | C_1) P(C_1).$$

Show that the Bayes error associated with this decision rule is

$$P_e = \frac{1}{\sqrt{2\pi}} \int_a^\infty e^{-z^2/2} dz$$

where $a = \frac{\mu_2 - \mu_1}{2\sigma}$.

4 Bias Variance for Ridge Regression

Recall the statistical model for ridge regression from lecture. We have a design matrix \mathbf{X} , where the rows of $\mathbf{X} \in \mathbb{R}^{n \times d}$ are our data points $\mathbf{x}_i \in \mathbb{R}^d$. We assume a linear regression model

$$Y = \mathbf{X}\mathbf{w}^* + \mathbf{z}$$

Where $\mathbf{w}^* \in \mathbb{R}^d$ is the true parameter we are trying to estimate, $\mathbf{z} = [z_1, \dots, z_n] \sim \mathcal{N}(0, I_n)$, and $Y = [y_1, \dots, y_n]$ is the random variable representing our labels.

Throughout this problem, you may assume $\mathbf{X}^\top \mathbf{X}$ is invertible. Given a realization of the labels $Y = \mathbf{y}$, recall these two estimators we have studied:

$$\begin{aligned}\mathbf{w}_{\text{ols}} &= \min_{\mathbf{w} \in \mathbb{R}^d} \|\mathbf{X}\mathbf{w} - \mathbf{y}\|_2^2 \\ \mathbf{w}_{\text{ridge}} &= \min_{\mathbf{w} \in \mathbb{R}^d} \|\mathbf{X}\mathbf{w} - \mathbf{y}\|_2^2 + \lambda \|\mathbf{w}\|_2^2\end{aligned}$$

- (a) Write the solution for \mathbf{w}_{ols} , $\mathbf{w}_{\text{ridge}}$. No need to derive it.
- (b) Let $\widehat{\mathbf{w}} \in \mathbb{R}^d$ denote any estimator of \mathbf{w}^* . In the context of this problem, an estimator $\widehat{\mathbf{w}} = \widehat{\mathbf{w}}(Y)$ is any function which takes the data \mathbf{X} and a realization of Y , and computes a guess of \mathbf{w}^* .

Define the MSE (mean squared error) of the estimator $\widehat{\mathbf{w}}$ as

$$\text{MSE}(\widehat{\mathbf{w}}) := \mathbb{E} \left[\left\| \widehat{\mathbf{w}} - \mathbf{w}^* \right\|_2^2 \right].$$

Above, the expectation is taken w.r.t. the randomness inherent in \mathbf{z} . Note that this is a multivariate generalization of the mean squared error we have seen previously.

Define $\widehat{\boldsymbol{\mu}} := \mathbb{E}[\widehat{\mathbf{w}}]$. Show that the MSE decomposes as such

$$\text{MSE}(\widehat{\mathbf{w}}) = \left\| \widehat{\boldsymbol{\mu}} - \mathbf{w}^* \right\|_2^2 + \text{Tr}(\text{Cov}(\widehat{\mathbf{w}})).$$

Note that this is a multivariate generalization of the bias-variance decomposition we have seen previously.

Hint: The inner product of two vectors is the trace of their outer product. Also, expectation and trace commute, so $\mathbb{E}[\text{Tr}(A)] = \text{Tr}(\mathbb{E}[A])$ for any square matrix A .

- (c) Show that

$$\mathbb{E}[\mathbf{w}_{\text{ols}}] = \mathbf{w}^*, \quad \mathbb{E}[\mathbf{w}_{\text{ridge}}] = (\mathbf{X}^\top \mathbf{X} + \lambda I_d)^{-1} \mathbf{X}^\top \mathbf{X} \mathbf{w}^*.$$

That is, \mathbf{w}_{ols} is an *unbiased* estimator of \mathbf{w}^* , whereas $\mathbf{w}_{\text{ridge}}$ is a *biased* estimator of \mathbf{w}^* .

- (d) Let $\gamma_1 \geq \gamma_2 \geq \dots \geq \gamma_d$ denote the d eigenvalues of the matrix $\mathbf{X}^\top \mathbf{X}$ arranged in non-increasing order. Show that

$$\text{Tr}(\text{Cov}(\mathbf{w}_{\text{ols}})) = \sigma^2 \sum_{i=1}^d \frac{1}{\gamma_i}, \quad \text{Tr}(\text{Cov}(\mathbf{w}_{\text{ridge}})) = \sigma^2 \sum_{i=1}^d \frac{\gamma_i}{(\gamma_i + \lambda)^2}.$$

Finally, use these formulas to conclude that

$$\text{Tr}(\text{Cov}(\mathbf{w}_{\text{ridge}})) < \text{Tr}(\text{Cov}(\mathbf{w}_{\text{ols}})) .$$

Hint: Remember the relationship between the trace and the eigenvalues of a matrix. Also, for the ridge variance, consider writing $\mathbf{X}^\top \mathbf{X}$ in terms of its eigen-decomposition $U \Sigma U^\top$.